



Digital Advertising

To auction “keywords”
or not to auction
that is the question
Gruppe 2



Agenda

- Einführung / «Business Case» (MAC)
- Data (PK)
 - Generierung
 - Normalisierung
- Code
 - Übersicht / Reward / Learning (RB)
 - Optimierung (Tensorboard / Optuna) (PZ)
- Analyse
 - Hyperparameter_tuning script (EO)
 - Tensorboard-analyzer / Visualization-ad_performance scripts (EO)
 - Data comparisons / Conclusions (EO)
- Lessons-Learned / Herausforderungen (all)



Einführung – Business Case (MAC)

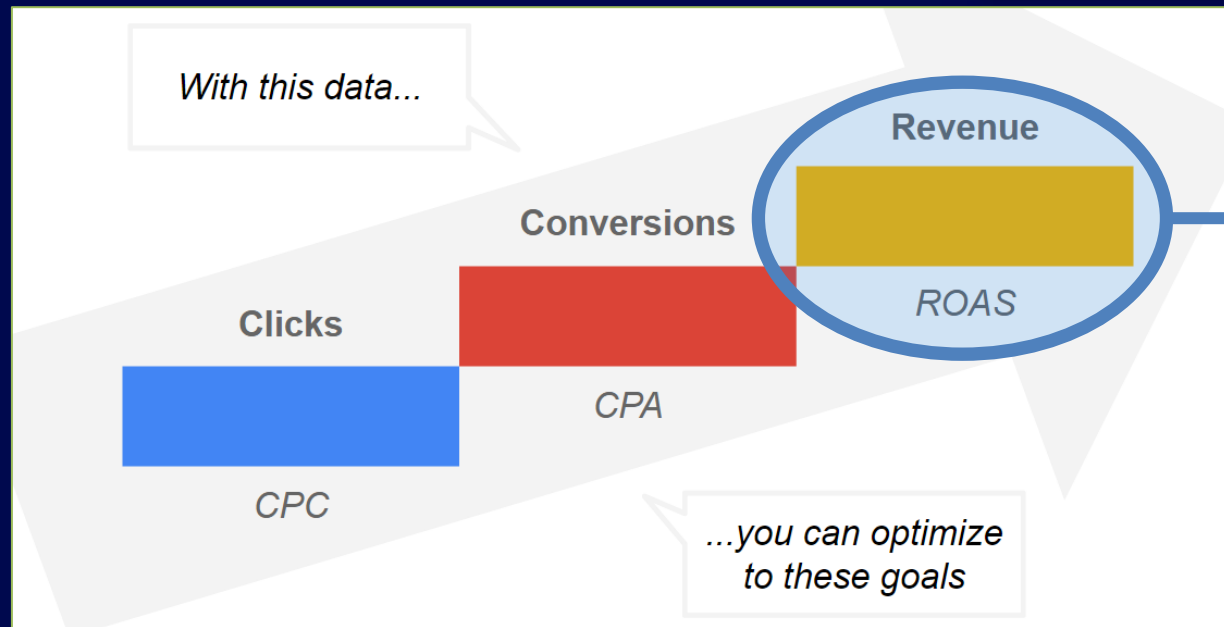
Einführung

- Datenbasierte Budgetallokation (10% des vorhandenen Bargeldes) dank Einsatz von KI bzw. Reinforcement Learning
- Von der Datenaufbereitung via Modell-Training zur fundierten Analyse und Verständnis

Business Case

- Unter der Annahme, dass (brauchbare) Daten vorliegen, rechnen wir mit einer Umsetzungszeit von 3-4 Monaten
- Danke einer Effektivitätssteigerung des ROAS um 20% erhalten wir einen positiven ROI nach 12 Monaten
- Mögliche Herausforderungen / Abhängigkeiten
 - Messbarkeit / Zuordnung der Effektivitätssteigerung
 - Produktions- / Lieferzeiten
 - Sonstiges (z.B. Landingpage-Qualität, andere Keywords ...)

Optimizing is all we need



Wir optimieren den Umsatz, der durch Werbeausgaben generiert wird, im Verhältnis zu den Kosten dieser Werbung.

Zustand (state (S_t))

- Vorhandenes «Bargeld» (cash)
- «Holdings» - aktive keywords

- keyword
- competitiveness
- difficulty_score
- organic_rank
- organic_clicks
- organic_ctr
- paid_clicks
- paid_ctr
- ad_spend
- ad_conversions
- ad_roas
- conversion_rate
- cost_per_click
- cost_per_acquisition
- previous_recommendation
- impression_share
- conversion_value

- done, terminated, truncated

Agent

Belohnung (reward (R_t))

- ROAS-basiert: ad_roas
→ maximiere Return on Ad Spend
- Kosteneffizienz: $ad_conversions / ad_spend$
→ maximieren Conversions pro Ausgabeeinheit
- Kombination: $0.6 * ad_roas + 0.4 * organic_clicks$
→ Vergangene Investition zahlt sich aus
- Nur «Organisches» Ranking: $organic_ctr$
→ Vergangene Investition zahlt sich aus

reward (R_{t+1})

Aktion (action (A_t))

- 1 Keyword kaufen/bieten oder nichts tun

action (A_{t+1})

Umgebung

state (S_{t+1})

Daten (PK)

- Suche nicht zielführend
 - [Shopping Mall Paid Search Campaign Dataset](#)
 - [PPC Campaign Performance Data](#)

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R
1	keyword	competitiveness	difficulty_score	organic_rank	organic_clicks	organic_ctr	paid_clicks	paid_ctr	ad_spend	ad_conversion	ad_roas	conversion_rate	cost_per_click	cost_per_acq	previous_reco	impression_share	conversion_value	
2	Autoversicherung	0.56	0.866	80	131	0.15703732	746	0.1274197	1350.26	39	1.43688913	0.05227882	1.81	34.6220513	1	0.27564632	1940.17391	
3	Kreditvergleich	0.72	0.309	95	34	0.01	465	0.09809572	1116	35	5.18030778	0.07526882	2.4	31.8857143	0	0.29563762	5781.22349	
4	Hotel buchen	0.34	0.806	75	117	0.07054571	107	0.03618353	131.61	4	4.39885891	0.03738318	1.23	32.9025	0	0.21776979	578.933821	
5	Flug buchen	0.8	0.214	34	308	0.21758421	485	0.09075523	1096.1	32	1.03442398	0.06597938	2.26	34.253125	1	0.5141958	1133.83212	
6	Mietwagen	0.88	0.349	86	59	0.01271818	1048	0.10847737	3028.72	57	0.71031707	0.05438931	2.89	53.1354386	0	0.44077783	2151.35151	
7	Online Bankin	0.43	0.413	23	157	0.13066824	402	0.12775712	590.94	16	4.93368431	0.039801	1.47	36.93375	0	0.24457658	2915.51141	
8	Sparplan	0.6	0.404	72	103	0.02719325	1202	0.13462178	372.62	66	22.8909404	0.05490849	0.31	5.64575758	0	0.31918163	8529.6222	
9	Reiseversiche	0.4	0.405	54	661	0.13385591	708	0.08440347	885	27	4.33914756	0.03813559	1.25	32.7777778	0	0.1	3840.14559	
10	Smartphone k	0.62	0.75	64	184	0.05797873	797	0.10085611	2255.51	33	2.35678364	0.04140527	2.83	68.3487879	0	0.23926687	5315.74906	
11	Laptop kaufen	0.66	0.615	32	51	0.08935746	203	0.10652401	491.26	13	4.32124211	0.06403941	2.42	37.7892308	0	0.21208059	2122.8534	
12	GÄstliche Ho	0.73	0.319	77	27	0.01	657	0.097323	361.35	40	14.1922005	0.0608828	0.55	9.03375	0	0.24584845	5128.35165	



Daten

Feature	In diesem Datensatz	Reiner <code>np.random.uniform</code>	Realitätsgrad
CTR	<code>np.random.beta(2, 5)</code>	Gleichverteilung	Mittel – Beta-Verteilungen modellieren oft Klickwahrscheinlichkeiten besser, da sie asymmetrisch sein können (z.B. viele niedrige CTRs, wenige hohe).
Bid	<code>np.random.uniform(0.1, 5.0)</code>	Gleichverteilung	Niedrig – realistischer wären Cluster (z.B. viele kleine Bids, wenige hohe).
Competition	<code>np.random.uniform(0.1, 1.0)</code>	Gleichverteilung	Niedrig – könnte auch von Keyword-Typ abhängen.
Paid Clicks	Binomial (Impressions, CTR)	-	Relativ gut – Binomial passt hier gut, da Klicks von Impressions & CTR abhängen.
Revenue	Zufällig pro Click	Gleichverteilung	Mittel – echte Einnahmen pro Klick hängen stark vom Keyword & Intent ab.
CPC	Bid x Competition	-	Okay – besser wäre ein dynamischeres Auktionsmodell.



Code –Dateien (RB)

- `digital_advertising.py` - Deep Reinforcement Learning Module
- `hyperparameter_tuning.py` - Hyperparameter tuning mit Optuna
- `visualize_ad_performance.py` - Leistungsvisualisierung
- `tensorboard-analyzer.py` - Trainingsanalyse
- `analyze_raw_data.py` - Interaktiver Rohdatenexplorer



Code – digital_advertising.py

- Daten laden und in Training und Test aufteilen
- Custom Environment erstellen
 - specs
 - _reset
 - _step
 - _compute_reward
- Policy erstellen
- Trainingsloop mit integrierter Verifikation mit Test und speichern des besten Models
- Vorhersagen (Inference)
- Speichern der Werte in Tensorboard



Code Custom Environment Specs

- Unsere Environment hat folgende Specs:

```
self.num_features = len(feature_columns)
self.num_keywords = get_entry_from_dataset(self.dataset, 0).shape[0]
self.action_spec = OneHot(n=self.num_keywords + 1) # select which one to buy or the last one to buy
nothing
self.reward_spec = Unbounded(shape=(1,), dtype=torch.float32)
self.observation_spec = Composite(
    observation = Composite(
        keyword_features=Unbounded(shape=(self.num_keywords, self.num_features), dtype=torch.float32),
        cash=Unbounded(shape=(1,), dtype=torch.float32),
        holdings=Bounded(low=0, high=1, shape=(self.num_keywords,), dtype=torch.int, domain="discrete")
    ),
    step_count=Unbounded(shape=(1,), dtype=torch.int64)
)
self.done_spec = Composite(
    done=Binary(shape=(1,), dtype=torch.bool),
    terminated=Binary(shape=(1,), dtype=torch.bool),
    truncated=Binary(shape=(1,), dtype=torch.bool)
)
```




Code Custom Environment - _step

- Holt die gewählte Action aus dem TensorDict
- Aktualisiert den State (Geld, welches Keyword wir kaufen)
- Berechnet den Reward
- Prüft, ob wir fertig sind
- Erhöht den aktuellen Schritt Index
- Aktualisiert TensorDict mit den aktuellen Werten und gibt den nächsten State gemäss TorchRL vorgaben zurück



Code Custom Environment – _step

- TorchRL: Aktueller State in SensorDict, neuer State als Rückgabewert

```
# tensordict is used from EnvBase later on, so we add the current state here
tensordict["done"] = torch.as_tensor(bool(terminated or truncated), dtype=torch.bool, device=self.device)
tensordict["observation"] = self.obs
tensordict["reward"] = torch.tensor(reward, dtype=torch.float32, device=self.device)
tensordict["step_count"] = torch.tensor(self.current_step-1, dtype=torch.int64, device=self.device)
tensordict["terminated"] = torch.tensor(bool(terminated), dtype=torch.bool, device=self.device)
tensordict["truncated"] = torch.tensor(bool(truncated), dtype=torch.bool, device=self.device)

# next as return value is also used by EnvBase and later added to tensordict by EnvBase
next_obs = TensorDict({
    "keyword_features": next_keyword_features, # next pki for each keyword
    "cash": cash_normalized.clone().detach(), # Current cash balance
    "holdings": self.holdings.clone()
}, batch_size=[])

next = TensorDict({
    "done": torch.tensor(bool(terminated or truncated), dtype=torch.bool, device=self.device),
    "observation": next_obs,
    "reward": torch.tensor(reward, dtype=torch.float32, device=self.device),
    "step_count": torch.tensor(self.current_step, dtype=torch.int64, device=self.device),
    "terminated": torch.tensor(bool(terminated), dtype=torch.bool, device=self.device),
    "truncated": torch.tensor(bool(truncated), dtype=torch.bool, device=self.device)
}, batch_size=tensordict.batch_size)

return next
```




Code – Custom Environment _compute_reward

- Ziel: Bester "Return on Ad Spend (ROAS)"

```
def _compute_reward(self, action, current_pki, action_idx, ad_roas):  
    """Compute reward based on the selected keyword's metrics"""  
    adjusted_reward = 0 if action_idx < self.num_keywords else 1 # encourage the agent to buy something  
    if ad_roas > 0: # log(0) is undefined  
        adjusted_reward = np.log(ad_roas) ## Adjust reward based on ad_roas performance, scale it with  
log  
    missing_rewards = []  
    # Calculate the ad_roas we did not get because we chose another keyword  
    for i in range(self.num_keywords):  
        sample = current_pki.iloc[i]  
        if action[i] == False:  
            missing_rewards.append(sample["ad_roas"])  
    # Adjust reward based on missing rewards to penalize the agent when not selecting keywords with  
    # high(er) ROAS  
    # clipping reduces the variance of the rewards  
    return np.clip(adjusted_reward - np.mean(missing_rewards) * 0.2, -2, 2)
```

- Wir brauchen die Policy für Training, Test und Inference

```
def create_policy(env, feature_dim, num_keywords, device):
    action_dim = env.action_spec.shape[-1]
    total_input_dim = feature_dim * num_keywords + 1 + num_keywords # features per keyword + cash + holdings

    # Create the flattening module to combine our observation into one vector
    flatten_module = TensorDictModule(
        FlattenInputs(),
        in_keys=[("observation", "keyword_features"), ("observation", "cash"), ("observation", "holdings")],
        out_keys=["flattened_input"]
    )

    value_mlp = MLP( # Create the value network
        in_features=total_input_dim,
        out_features=action_dim,
        num_cells=[256, 256, 128, 64], # Deeper and wider architecture
        activation_class=nn.ReLU # ReLU often performs better than Tanh
    )

    value_net = TensorDictModule(value_mlp, in_keys=["flattened_input"], out_keys=["action_value"])

    # Combine into the complete policy
    policy = TensorDictSequential(flatten_module, value_net, QValueModule(spec=env.action_spec))

    return policy.to(device)
```




Code Training Loop

- Normaler TorchRL Trainingsloop


```
exploration_module = EGreedyModule(  
    env.action_spec, annealing_num_steps=100_000, eps_init=exploration_eps_init, eps_end=exploration_eps_end  
)  
exploration_module = exploration_module.to(device)  
policy_explore = TensorDictSequential(policy, exploration_module).to(device)  
  
collector = SyncDataCollector(  
    env,  
    policy_explore,  
    frames_per_batch=frames_per_batch,  
    total_frames=-1,  
    init_random_frames=init_rand_steps,  
)  
replay_buffer_size = 100_000  
rb = ReplayBuffer(storage=LazyTensorStorage(replay_buffer_size))  
  
loss = DQNLoss(value_network=policy, action_space=env.action_spec, delay_value=True).to(device)  
  
optim = Adam(loss.parameters(), lr=lr, weight_decay=weight_decay) # Add weight decay for regularization  
updater = SoftUpdate(loss, eps=softupdate_eps)
```

Optuna

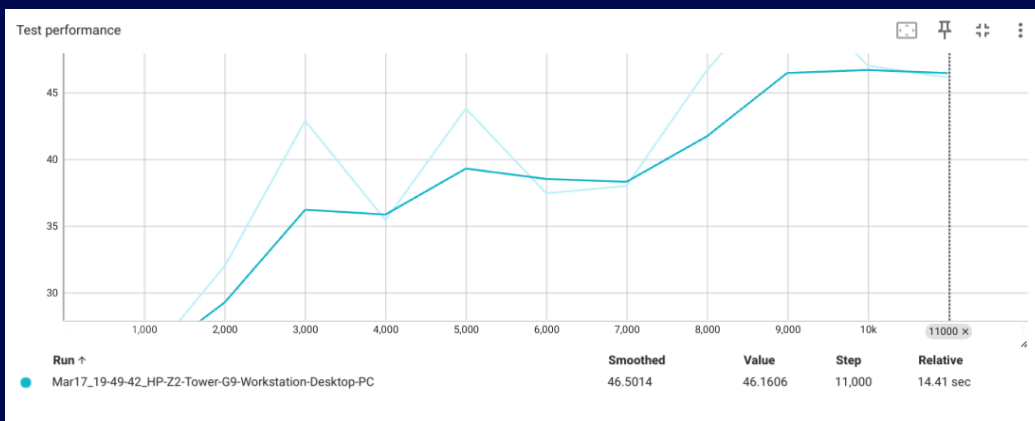
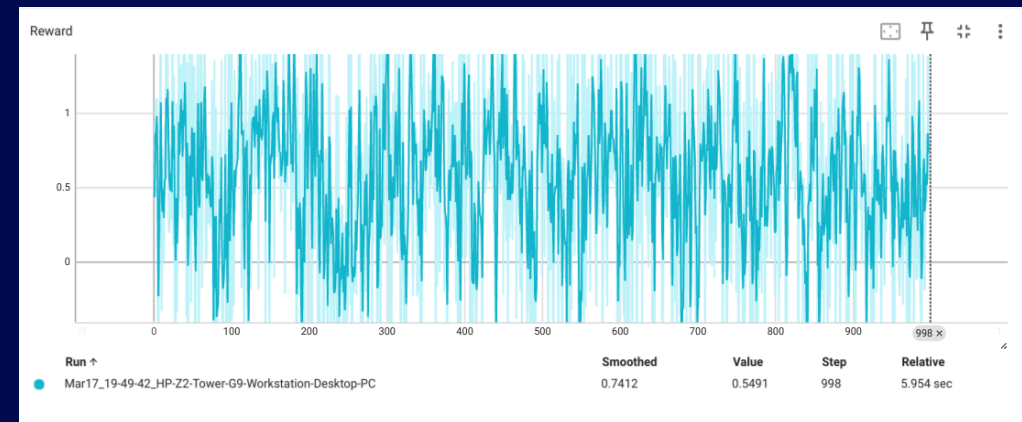
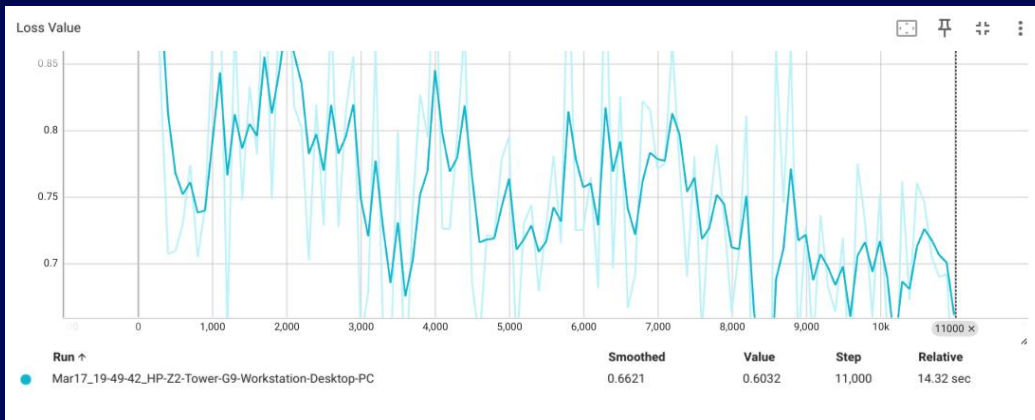
```
# Learning rate for the optimizer
# Batch size for training
), # Initial value for epsilon in epsilon-greedy exploration
, # Final value for epsilon in epsilon-greedy exploration
# Soft update rate for target network
# Discount factor for future rewards
# Weight decay for regularization
```

```
# Sample hyperparameters
params = {
    'lr': trial.suggest_float('lr', 1e-4, 1e-2, log=True),
    'batch_size': trial.suggest_categorical('batch_size', [32, 64, 128, 256]),
    'exploration_eps_init': trial.suggest_float('exploration_eps_init', 0.5, 1.0),
    'exploration_eps_end': trial.suggest_float('exploration_eps_end', 0.01, 0.1),
    'softupdate_eps': trial.suggest_float('softupdate_eps', 0.9, 0.99),
    'gamma': trial.suggest_float('gamma', 0.9, 0.99),
    'weight_decay': trial.suggest_float('weight_decay', 1e-6, 1e-4)
}
```

Hyperparameter Importance

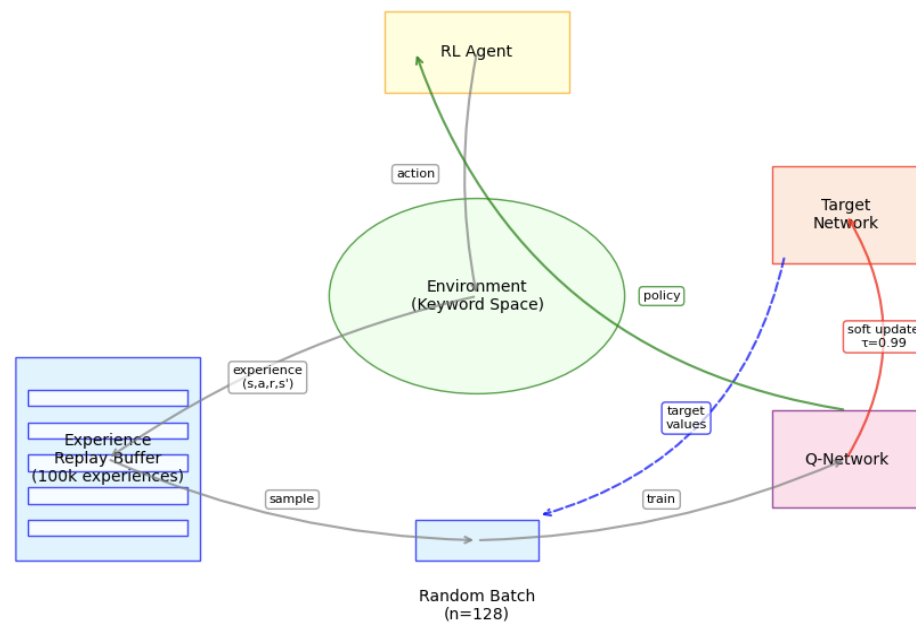


Tensorboard




Data analysis (EO)

Experience Replay and Target Network Architecture




- Experience Replay Buffer: Stores 100,000 transitions (s,a,r,s') to break temporal correlations
- Random Sampling: Draws minibatches of 128 experiences for decorrelated updates
- Target Network: Stabilizes learning by providing consistent targets during updates
- Soft Updates: Gradually transfers Q-network weights to target network ($\tau=0.99$)

Impact of Experience Replay and Target Network on Training Stability

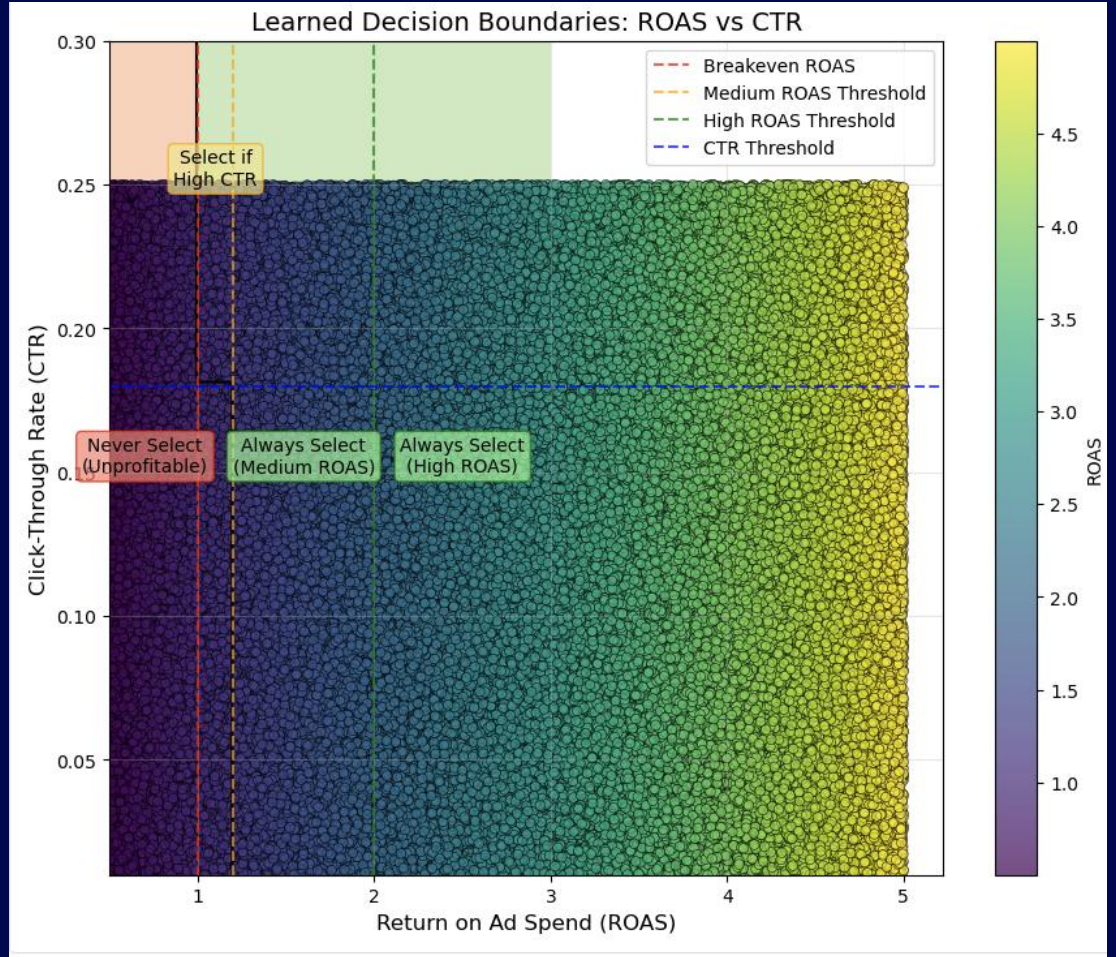
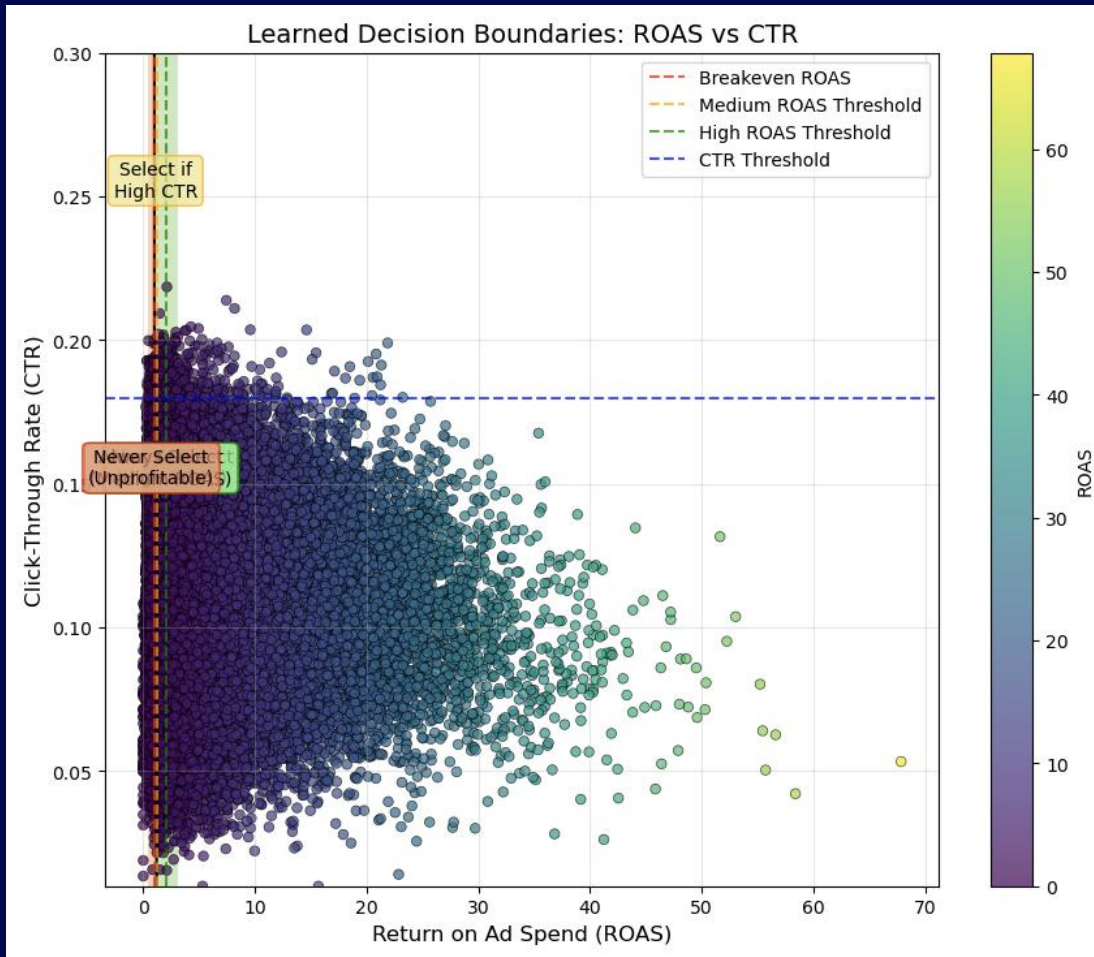




Cash Management Overview

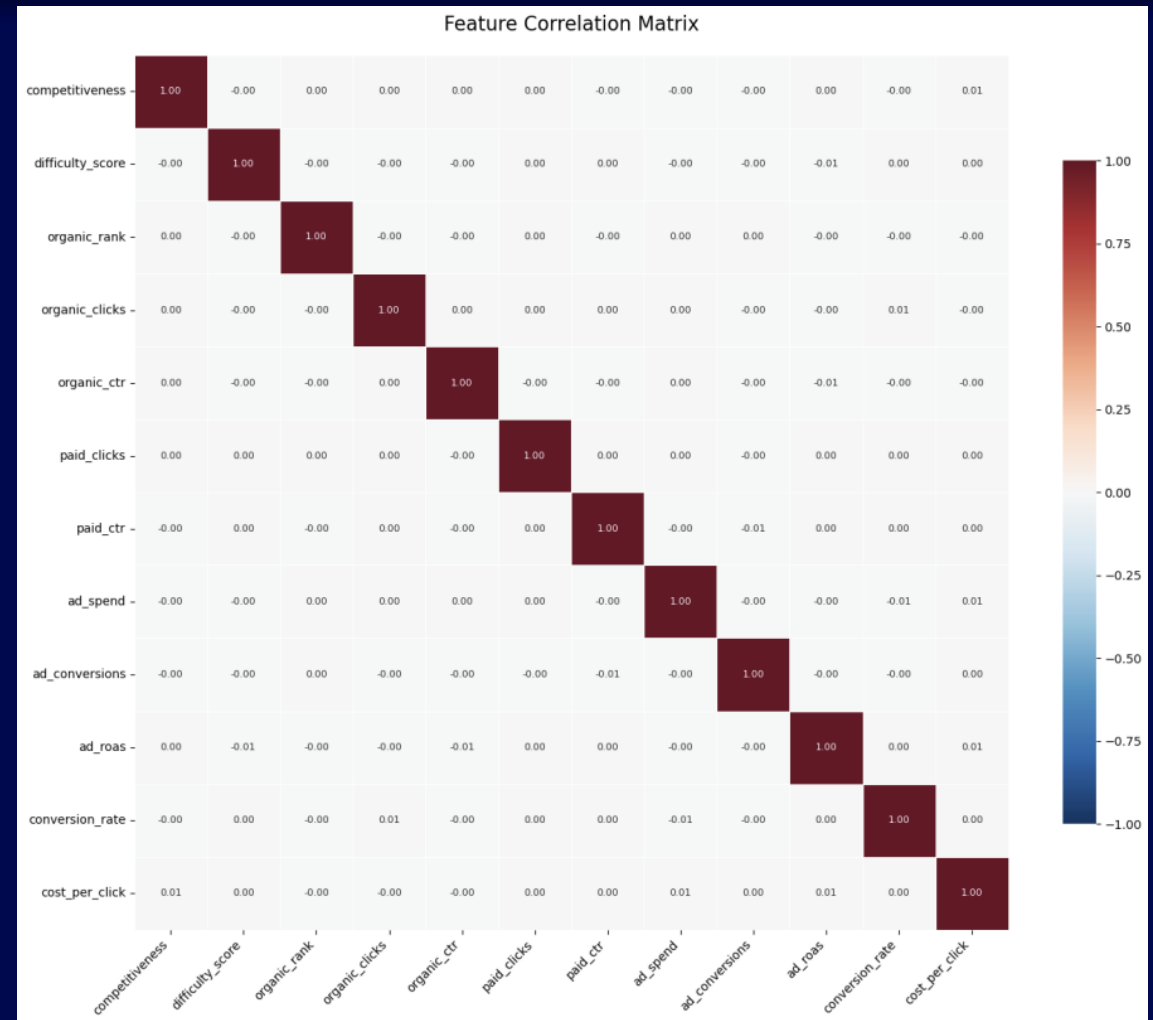
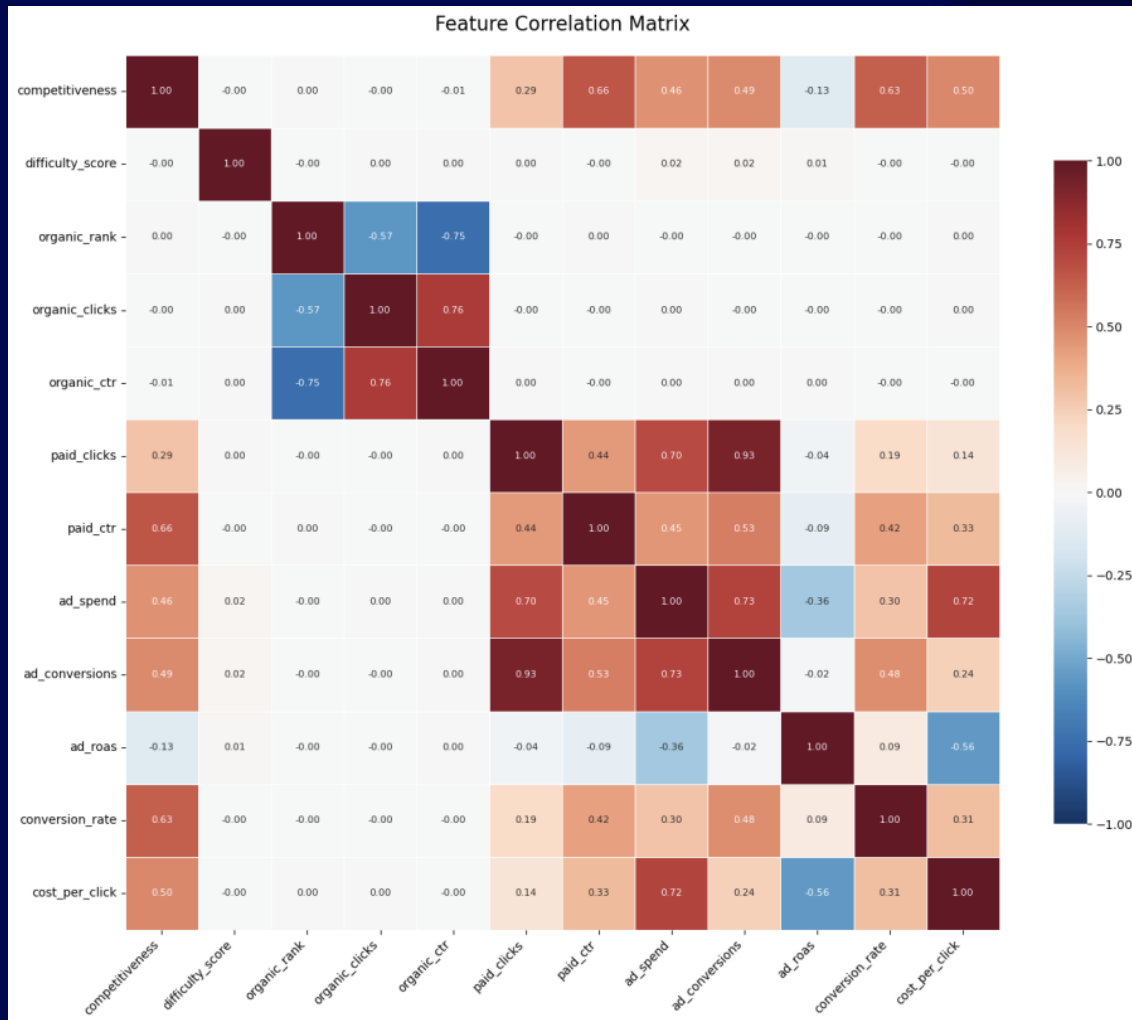


Comparison: Our data vs "Ilja simple data"


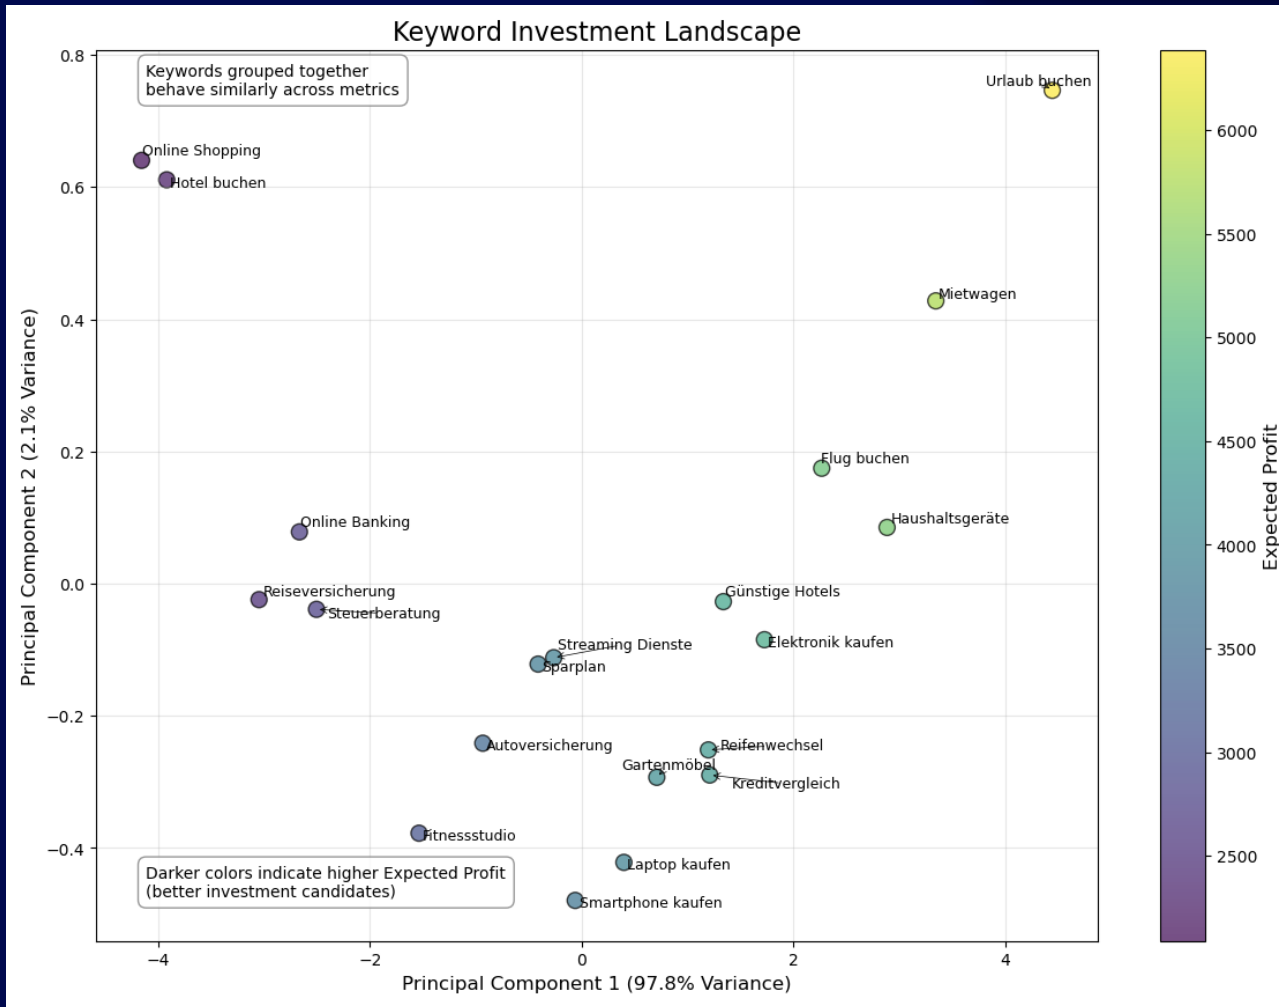




Comparison: Our data vs "Ilja simple data"



Which keywords to invest in?





Lessons Learned / Herausforderungen (MAC)

- Erfahrungen in einem „realen“ Business-Projekt
→ von der Idee via Daten und Training zur Analyse
- Daten: Gewisse offene Fragen hinsichtlich Datenqualität
- TorchRL Library
- Fehlendes Knowhow über "Digital Advertising" aufgebaut



GitHub- / Colab-Link

- [Digital Advertising «Gruppe 2» on GitHub](#)
- [Colab-Notebook](#)